

Bachelor Thesis Midterm

A Multimodal XR Museum with vitrivr and Godot

Jannick Seper

University of Basel

22 July 2026

VIRTUE no longer matches today's media and XR landscape

01

Aged technology

The predecessor relies on Cineast and Cottontail DB.

02

Limited media scope

VIRTUE primarily supports images rather than heterogeneous collections.

03

New XR landscape

Modern headsets provide new interaction modes across multiple platforms.

A modern XR museum for similarity-based exploration of multimodal collections.

VIRTUE into a multimodal XR museum

PREDECESSOR

VIRTUE

- › dynamically arranged image galleries in VR
- › retrieval with spatial exploration
- › Cineast and Cottontail DB
- › image collections



THIS THESIS

Multimodal XR museum

- › vitrivr-engine and pgvector
- › images, videos, and 3D models
- › multiple headsets through OpenXR
- › similarity-driven exploration

Background: key terms

Term	Meaning	Role in this project
VR	Virtual Reality	The fully virtual museum environment experienced through a headset.
XR	Extended Reality	The umbrella term covering VR, augmented reality, and mixed reality.
OpenXR	Cross-platform XR API	A common interface for headsets, controllers, tracking, and XR runtimes.
CLIP	Contrastive Language–Image Pre-training	Maps text and images to comparable numerical embeddings.
HUD	Head-up display	Shows search progress and status messages within the user's view.

Three goals define the prototype

01

Retrieve

vitivr-engine and pgvector provide semantic search over multimodal embeddings.

02

Experience


The XR application presents retrieved images, videos, and 3D models as interactive exhibits.

03

Explore

Users trigger new similarity searches from exhibits.

Continuous exploration loop




One concept

The interaction is independent of whether a result is an image, video, or 3D model.

Two entry points


Exploration can start from language or continue directly from an existing exhibit.

CLIP creates a shared comparison space




Key distinction: the collection is encoded once; a live search only encodes the new query and compares it with the stored vectors.

Rendered views make 3D models searchable



Rather than encoding the mesh itself, the system searches over rendered views. A matching view points back to the original GLB model, which is then loaded into the museum.


Three responsibilities shape the system



Design decision: separate experience, orchestration, and retrieval so each can evolve independently.


vitivr request supports both search modes

SIMPLIFIED REQUEST-RESPONSE FLOW



Only the input mapping changes. Response parsing, media loading, and VR placement are shared. The returned vector enables the next result-to-result search.

Godot and vitrivr remain replaceable



Why Godot?

01

Godot XR Tools

Ready-made building blocks for movement, interaction, grabbing, and XR interfaces.


02

OpenXR Vendors

Simple device integration for Quest and Focus; Windows uses the regular OpenXR runtime.

Controller and hand input use the same scenes and application logic.

Entering the museum




While in the menu

The museum is hidden, player movement is disabled, and configuration happens inside XR.

On transition

`PlatformSwitcher` moves the XR rig, swaps environment and HUD, and preserves museum state.

Media become exhibits



2D MEDIA

Images and videos use centred wall slots and aspect-aware frames.

3D MODELS

GLB objects are instanced on stands and fitted to their bounds.

Every exhibit stores its path, name, and CLIP vector for **Similar**. 3D models also support **Original Size**.

Demo

Search → explore → inspect → search again

Evaluation questionnaire

01 SUS

10 standard statements about usability.

02 Project questions

Navigation, controls, responsiveness, 3D interaction, and design.

03 Context + feedback

VR experience, device, usage mode, and one open question.

Next steps

1. Resolve slow media loading issues.
2. Integrate and test streaming to Apple Vision Pro.
3. Run user evaluations.
4. Complete the reproducible setup and documentation.
5. Finish the thesis.